

[Back to Week 7](#)[× Lessons](#)

This Course: CS 498: Cloud Networking

[Prev](#)[Next](#)

## Programming Assignment: Unicast Routing

You have not submitted. You must earn 48/80 points to pass.

**Deadline** Pass this assignment by November 4, 09:59 PM PST

**Instructions** [My submission](#)

## Programming Assignment 2: Unicast Routing

In this assignment, you will implement traditional shortest-path routing, with the

How to submit  
When you're ready to submit, you can upload files

link state (LS) or distance/path vector (DV/PV) protocols.

The test setup will be the same environment as in MP1. Your nodes will use different virtual network adapters, with iptables restricting communication among them. We are providing you with the same script that our testing environment uses to establish these restrictions.

This MP introduces the unfortunate distinction between “network order” and “host order”, AKA endianness. The logical way to interpret 32 bits as an integer is to call the left-most bit the highest, and the right-most the lowest. This is the standard that programs sending raw binary integers over the internet are expected to follow. Due to historical quirks, however, x86 processors store integers reversed at the byte level. Look up `htons()`, `htonl()`, `ntohs()`, `ntohl()` for information on how to deal with it.

for each part of the assignment on the "My submission" tab.

programmingassignment2\_files.zip

## Routing Environment:

For this assignment, a virtual network will be emulated using the iptables in Linux. iptables is an application that allows users to configure specific rules that will be enforced by the kernel's netfilter framework. It acts as a packet filter and firewall that examines and directs traffic based on port, protocol and other criteria.

Your nodes will all run on the same machine. There will be a made-up topology applied to them in the following manner:

- For each node, a virtual interface (eth0:1, eth0:2, etc) will be created and given an IP address.
- A node with ID  $n$  gets address 10.1.1. $n$ . IDs 0-255 inclusive are valid.

- Your program will be given its ID on the command line, and when binding a socket to receive UDP packets, it should specify the corresponding IP address (rather than INADDR\_ANY / NULL).
- iptables rules will be applied to restrict which of these addresses can talk to which others. 10.1.1.30 and 10.1.1.0 can talk to each other if and only if they are neighbors in the made-up topology.
- The topology's connectivity is defined in a file that gets read by the script that creates the environment. The link costs are defined in files that tell nodes what their initial link costs to other nodes are, **if they are in fact neighbors.**
- A node's only way to determine who its neighbors are is to see who it can directly communicate with. Nodes do not get to see the aforementioned connectivity file.

## Manager:

While running, your nodes will receive instructions and updated information from a manager program. You are not responsible for submitting an implementation of this manager; we will test using our own. Your interaction with the manager is very simple: it sends messages in individual UDP packets to your nodes on UDP port 7777, and they do not reply in any way.

The manager's packets have one of two meanings: "send a data packet into the network", or "your direct link to node X now has cost N." Their formats are:

### **"Send a data packet":**

**"send"**<4 ASCII bytes>**destID**<net order 2 bytes>**message**<some ASCII message (shorter than 100 bytes)>

Example: destination Id is 4 and message is "hello". The manager

command will be:

`"send4hello"`

where 4 occupies 2 bytes and is in the network order. You need to convert it to the host order to get the correct destID. Note that there is no space delimiter among "send", destID and the message body.

**"New cost to your neighbor":**

**"cost"**<4 ASCII bytes>**destID**<net order 2 bytes>**newCost**<net order 4 bytes>

Example: the manager informs node 2 that the cost to node 5 is set to 33. The command received by node 2 will be:

`"cost533"`

where 5 occupies 2 bytes and 33 occupies 4 bytes. Both of them are in the network order. You need to convert them to the host order to get the correct destID and newCost. Note that

there is no space  
delimiter among "cost",  
destID and newCost.

The source code of the  
manager is included in  
the attachment. You  
can read it to gain a  
better understanding  
of the format. Feel free  
to compile it and use it  
to test your nodes.

## Your Nodes:

Whether you are  
writing an LS or DV/PV  
node, your node's  
interface to the  
assignment  
environment is the  
same. Your node  
should run like:

```
./ls_router nodeid  
initialcostsfile logfile
```

```
./vec_router nodeid  
initialcostsfile logfile
```

Examples:

```
./ls_router 5  
node5costs logout5.txt
```

```
./vec_router 0 costs0.txt  
test3log0
```

When originating,  
forwarding, or receiving  
a data packet, your  
node should log the

event to its log file. The sender of a packet should also log when it learns that the packet was undeliverable. The format of the logging is described in the next section. **Note:** even if the node has nothing to write to the log file, the log file should still be created.

The link costs that your node receives from the input file and the manager don't tell your node whether those links actually currently exist, just what they would cost if they did. Your node therefore needs to constantly monitor which other nodes it is able to send/receive UDP packets directly to/from. In the attachment (main.c and monitor\_neighbors.h), we have provided the code that you can use for this.

That concludes the description of how your nodes need to do I/O, interact with the manager program, and stay aware of the



topology. Now, for what they should actually accomplish:

- Using LS or DV/PV, maintain a correct forwarding table.
- Forward any data packets that come along according to the forwarding table.
- React to changes in the topology (changes in cost and/or connectivity).
- Your nodes should converge within 5 seconds of the most recent change.

**Partition:** the network might become partitioned, and your protocols should react correctly: when a node is asked to originate a packet towards a destination it does not know a path for, it should drop the packet, and rather than log a send event, log an unreachable event (see File Formats section).

**Tie breaking:** we would like everyone to have consistent output even on complex

topologies, so we ask you to follow specific tie-breaking rules.

- DV/PV: when two equally good paths are available, your node should choose the one whose next-hop node ID is lower.

- LS: When choosing which node to move to the finished set next, if there is a tie, choose the lowest node ID.

## File Formats

Your nodes take the “initial costs file” as input, and write their output to a log file. The locations of both of these files are specified on the command line, as described earlier. The initial costs files might be read-only.

### **Initial costs file format:**

<nodeid> <nodecost>

<nodeid> <nodecost>

### **Example initial costs file:**

5 23453245

2 1

3 23

19 1919

200 23555

In this example, if this file was given to node 6, then the link between nodes 5 and 6 has cost 23453245 – as long as that link is up in the physical topology.

If you don't find an entry for a node, default to cost 1. We will only use positive costs – never 0 or negative. We will not try to break your program with malformed inputs. Once again, just because this file contains a line for node n, it does NOT imply that your node will be neighbors with n.

**Log file:** (See our provided code for sprintf()s that generate these lines correctly.)

**Example log file:**

forward packet dest  
[nodeid] nexthop  
[nodeid] message [text  
text]

sending packet dest  
[nodeid] nexthop  
[nodeid] message [text  
text]

receive packet message  
[text text text]

unreachable dest  
[nodeid]

...

In this example, the node forwarded a message bound for node 56, received a message for itself, originated packets for nodes 11 and 12 (but realized it couldn't reach 12), then forwarded another two packets.

forward packet dest 56  
nexthop 11 message  
Message1

receive packet message  
Message2!

sending packet dest 11  
nexthop 11 message  
hello there!

unreachable dest 12

forward packet dest 23  
nexthop 11 message  
Message4

forward packet dest 56  
nexthop 11 message  
Message5

Our tests will have data  
packets be sent  
relatively far apart, so  
don't worry about  
ordering.

## Extra Notes:

To set an initial  
topology, run perl  
make\_topology.pl  
thetopofile (an example  
topology file is  
provided). Note that  
you need to replace all  
the "eth0" in the  
make\_topology.pl file to  
your VM's network  
name (e.g., enp0s3) to  
make the iptables  
work. This will set the  
connectivity between  
nodes. To give them  
their initial costs,  
provide each with an  
initialcosts file  
(examples of these are  
provided as well. See  
the spec document for  
more details on all file  
formats). You can just  
give them all an empty  
file, if you want all links  
to default to cost 1.

To make the nodes send messages, or change link costs while the programs are running, use `manager_send`.

To bring links up and down while running:

e.g. to bring the link between nodes 1 and 2 up:

```
sudo iptables -I  
OUTPUT -s 10.1.1.1 -d  
10.1.1.2 -j ACCEPT ;  
sudo iptables -I  
OUTPUT -s 10.1.1.2 -d  
10.1.1.1 -j ACCEPT
```

To take that link down:

```
sudo iptables -D  
OUTPUT -s 10.1.1.1 -d  
10.1.1.2 -j ACCEPT ;  
sudo iptables -D  
OUTPUT -s 10.1.1.2 -d  
10.1.1.1 -j ACCEPT
```

log file format:

```
sprintf(logLine,  
"sending packet dest  
%d nexthop %d  
message %s\n", dest,  
nexthop, message);
```

```
sprintf(logLine,  
"forward packet dest  
%d nexthop %d
```

```
message %s\n", dest,  
nexthop, message);
```

```
sprintf(logLine, "receive  
packet message %s\n",  
message);
```

```
sprintf(logLine,  
"unreachable dest  
%d\n", dest);
```

```
... and then  
fwrite(logLine, 1,  
strlen(logLine),  
theLogFile);
```

If that's all you do to the log file, you will have the correct format. (Don't worry about closing it; when your process gets killed, any output should be flushed + the file will be closed).

You can use the provided files as a starting point for the routing protocol implementation. However, it is fine if you want to start from scratch.

## Grading

The grading will be determined by 8 tests, which are supposed to

be progressively harder and stress different elements of your code.

- Send a message to a non-neighbor node in a 3 node topology
- Send a message (by shortest path, of course) to a non-neighbor in a certain small (<20 node) topology.
- Send a message to a non-neighbor node in a certain large (>50 node) topology
- Switch to a better path when one becomes available in a certain large topology
- Correctly fall back to a worse path on a certain small topology
- Correctly report a failed send in a certain small topology that gets partitioned
- Get a packet through after a former partition is healed, large topology
- Converge within the time limit after a major, rapid series of changes



to the network (The time limit starts when the final change is made).

### **Grading Script:**

Run the following grading script according to the instructions provided along with it and submit your results.

MP2\_gradingscript.zip

-Copy your compressed code files (compressed into "code.zip") into the grading script directory. There is no format requirement for the code files. We keep the code for plagiarism check.

-Copy your compiled program in the grading script directory.

-Figure out the name of the network the router will run on. For the routing to work smoothly make sure your VM is connected to only one network. You can use the "sudo ifconfig" command to see the name of the network.

-Run the grading script as follows:

```
"python autograde.py  
[program_name]  
[network_name]"
```

Example: if  
program\_name is  
"vec\_router" and the  
network\_name is  
"enp0s3", the  
command will be:

```
"python autograde.py  
vec_router enp0s3"
```

-After the successful execution of the grading script, a zip file named out.zip will be created with your program results. Submit this file on the assignment submission page on coursera.

**NOTE:** Before running the grading script, make sure your program works correctly with the example topology provided in the programming assignment files.



